

RESEARCH

Finding Individual Strategies for Storage Units in Electricity Market Models using Deep Reinforcement Learning

Nick Harder^{1*}, Anke Weidlich¹ and Philipp Staudt²

*Correspondence:

nick.harder@inatech.uni-freiburg.de

¹Institute for Sustainable Systems Engineering, University of Freiburg, Freiburg im Breisgau, Germany

Full list of author information is available at the end of the article

Abstract

Modeling energy storage units realistically is challenging as their decision-making is not governed by a marginal cost pricing strategy but relies on expected electricity prices. Existing electricity market models often use centralized rule-based bidding or global optimization approaches, which may not accurately capture the competitive behavior of market participants. To address this issue, we present a novel method using multi-agent deep reinforcement learning to model individual strategies in electricity market models. We demonstrate the practical applicability of our approach using a detailed model of the German wholesale electricity market with a complete fleet of pumped hydro energy storage units represented as learning agents. We compare the results to widely used modeling approaches and demonstrate that the proposed method performs well and can accurately represent the competitive behavior of market participants. To understand the benefits of using reinforcement learning, we analyze overall profits, aggregated dispatch, and individual behavior of energy storage units. The proposed method can improve the accuracy and realism of electricity market modeling and help policymakers make informed decisions for future market designs and policies.

Keywords: agent-based modeling; electricity markets; energy storage; multi-agent reinforcement learning; reinforcement learning

Supplementary Material

Algorithm 1 MATD3 algorithm for bidding of N energy storage units in an electricity market

```

1: Initialize critic  $\theta_{1,2}$  and actor  $\phi$  for each learning agent
2: Initialize target critic  $\theta'_{1,2}$  and actor  $\phi'$  for each learning agent
3: for episode  $\in [1, M]$  do
4:   Initialize market simulation
5:   for  $t \in [1, T]$  do
6:     Create a global observation  $o_t^{\text{global}}$ 
7:     for  $i \in [1, N]$  do
8:       Derive individual observation  $o_{t,i}$  using the global and local observations
9:       Perform an action  $a_{t,i} = \pi_{\phi_i}(o_{t,i}) + \epsilon$ 
10:      Derive the bid  $B_{t,i} = (P_{t,i}, ep_{t,i})$  using the action  $a_{t,i}$ 
11:    end for
12:    Perform the market clearing
13:    for  $i \in [1, N]$  do
14:      Get accepted capacity  $P_{t,i}^{\text{conf.}}$  and market clearing price  $M_t$ 
15:      Calculate individual reward  $R_{t,i}$ 
16:      Store transition  $(s_{t,i}, a_{t,i}, r_{t,i}, s_{t+1,i})$  into the replay buffer
17:    end for
18:    if remainder  $\frac{t}{\text{train.freq.}}$  is 0 then
19:      for step  $\in [1, \text{grad.steps}]$  do
20:        for  $i \in [1, N]$  do
21:          Sample a minibatch of  $k \in B$  samples from the replay buffer
22:          Calculate target actions:  $\bar{\mathcal{A}}_{k,i} = \pi_{\phi'_i}(o'_{k,i}) + \epsilon$ 
23:          Calculate the target value:  $y_{k,i} = r_{k,i} + \gamma \min_{j=1,2} Q_{\theta'_{i,j}}(o'_{k,i}, \bar{\mathcal{A}}_{k,i})$ 
24:          Update the critics using gradient ascent:  $L(\theta_{i,j}) = \frac{1}{B} \sum_{k=1}^B [y_{k,i} - Q_{\theta_{j,i}}(o_{k,i}, \mathcal{A}_{k,i})]^2$ 
25:        if remainder  $\frac{\text{step}}{\text{delay}}$  is 0 then
26:          Update the actor:  $\nabla_{\phi} J(\phi_i) = \frac{1}{B} \sum_{k=1}^B \nabla_{\phi} \pi_{\phi_i}(o_{k,i}) \nabla_a Q_{\theta_{i,1}} \left( o_{k,i}, \mathcal{A}_{k,i}^{\setminus [i]}, \pi_{\phi_i}(o_{k,i}) \right)$ 
27:          Perform soft-updates:
28:           $\theta'_{i,j} = (1 - \tau)\theta'_{i,j} + \tau\theta_{i,j}$  for  $j = 1, 2$ 
29:           $\phi'_i = (1 - \tau)\phi'_i + \tau\phi_i$ 
30:        end if
31:      end for
32:    end if
33:  end for
34: end for
35: end for

```

Table 1: Actor and critic NN parameters and learning hyper-parameters.

Parameter	Value	Parameter	Value
Critic NN architecture	MLP, (\times _input, 512, 256, 128, 1)	Training frequency	1000
Actor NN architecture	MLP, (\times _input, 256, 128, 1)	Gradient steps	1000
Critic activations	(ReLU, ReLU, ReLU, Linear)	Buffer size	$5 \cdot 10^5$
Actor activations	(ReLU, ReLU, Tanh)	Policy delay, d	2
Observation size	128	Soft-update, τ	0.05
Action size	2	Target noise σ	0.2
Optimizer, learning rate	Adam, 10^{-4}	Target noise clip, c	0.5
Batch size, B	256	Learning starts	5
Reward discount, γ	0.999	Action noise	Gaussian

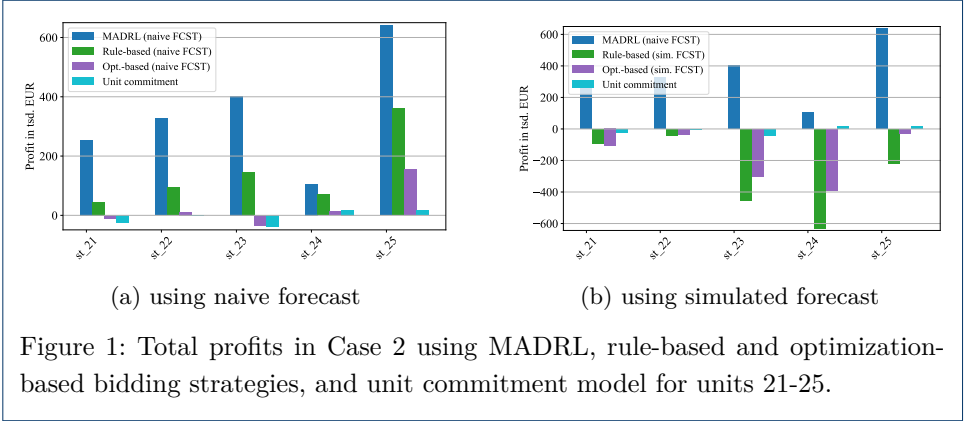


Figure 1: Total profits in Case 2 using MADRL, rule-based and optimization-based bidding strategies, and unit commitment model for units 21-25.